

Design patterns:

- **Factory:** Data returned by the application back to the client (or between the services) might have to be in the future in different formats than JSON (for example XML), so using a factory makes it easy to add new formats in the future.
- **Strategy:** Depending on the type of user, different methods and implementation will be available for them.

Issues:

Must haves:

- User
 - Users have to be authorised when trying to make a request
 - Users can make job requests
 - Assign user to single/multiple faculties
 - Users can have multiple outstanding requests
 - Users must be stored in the database
- Job
 - Users must be able to create new jobs
 - SysAdmins must be able to request the overview of all schedules
 - FacultyAccounts must be able to approve jobs
 - Job service must provide a job instance for the Scheduler service
 - Job service must be able to update the schedule of the jobs
 - Job service must be able to provide their status to the owner of the job
 - Jobs should be stored in the database (so that they can be accessed by the SysAdmin)
- Schedule
 - Job microservice must be able to request a job to be scheduled
 - Schedule must be able to determine whether or not a given job can be scheduled.
 - Schedule assigns a job to a particular faculty or rejects it based (or to the free pool, or both) on the resources and date preferred by
- Clusters
 - Users must be able to add a new node
 - Users must be able to remove a node
 - FacultyAccounts must be able to free the cluster's resources for specified days
 - Clusters must contain information about the particular faculty (free pool)

Should haves:

- User
 - Resign user of faculty
- Job
 - Users should be able to remove their jobs
- Schedule
 - Job microservice should be able to unschedule a job
- Clusters

Could have:

- User
- Job
- Schedule
 - Scheduler could be able to reschedule jobs greedily when new ones come in
- Clusters